

CoffeeSalesSyntheticDataPreparation

August 13, 2025

```
[1]: import pandas as pd
import numpy as np
import random

# Set random seed for reproducibility
np.random.seed(42)

# Parameters
dates = pd.date_range(start="2024-01-01", end="2024-03-31", freq="D")
locations = ["Tbilisi", "Batumi", "Kutaisi", "Rustavi"]
products = ["Espresso", "Latte", "Cappuccino", "Americano", "Mocha", "Cold_Brew"]
campaigns = ["Facebook Ads", "Google Ads", "Instagram Ads", "Billboard", "Flyers"]

# Generate sales data
sales_data = []
for date in dates:
    for loc in locations:
        for prod in products:
            units_sold = np.random.poisson(lam=10)
            price = random.choice([2.5, 3.0, 3.5, 4.0])
            revenue = round(units_sold * price, 2)
            sales_data.append([date, loc, prod, units_sold, price, revenue])

sales_df = pd.DataFrame(sales_data, columns=["Date", "Location", "Product", "Units_Sold", "Unit_Price", "Revenue"])

# Generate marketing data
ads_data = []
for date in dates:
    for camp in campaigns:
        clicks = np.random.poisson(lam=20)
        impressions = clicks * random.randint(20, 50)
        spend = round(clicks * random.uniform(0.1, 0.5), 2)
        ctr = round(clicks / impressions, 4) if impressions > 0 else 0
        cpc = round(spend / clicks, 2) if clicks > 0 else 0
```

```

        ads_data.append([date, camp, clicks, impressions, spend, ctr, cpc])

ads_df = pd.DataFrame(ads_data, columns=["Date", "Campaign", "Clicks", "Impressions", "Spend", "CTR", "CPC"])

# Save datasets
coffee_sales_path = "coffee_sales_data.csv"
coffee_ads_path = "coffee_ads_data.csv"

sales_df.to_csv(coffee_sales_path, index=False)
ads_df.to_csv(coffee_ads_path, index=False)

```

```

[2]: # ----- 1) Load -----
sales_df = pd.read_csv(coffee_sales_path)
ads_df = pd.read_csv(coffee_ads_path)

```

```

[2]: (
      Date Location      Product  Units_Sold  Unit_Price  Revenue
2179 2024-03-31  Rustavi      Latte           11          2.5     27.5
2180 2024-03-31  Rustavi  Cappuccino           13          4.0     52.0
2181 2024-03-31  Rustavi  Americano            4          3.0     12.0
2182 2024-03-31  Rustavi      Mocha            6          4.0     24.0
2183 2024-03-31  Rustavi  Cold Brew           16          3.5     56.0,
      Date      Campaign Clicks  Impressions  Spend      CTR      CPC
450 2024-03-31  Facebook Ads      17          561    3.29  0.0303  0.19
451 2024-03-31   Google Ads      20         1000    4.11  0.0200  0.21
452 2024-03-31 Instagram Ads      22          858    8.09  0.0256  0.37
453 2024-03-31   Billboard      21          945    4.05  0.0222  0.19
454 2024-03-31      Flyers      14          336    6.63  0.0417  0.47)

```

```

[3]: # ----- 2) Quick inspect -----
print("\n=== Raw shapes ===")
print("sales_df:", sales_df.shape, "ads_df:", ads_df.shape)

print("\n=== Sales head ===")
print(sales_df.head())

print("\n=== Ads head ===")
print(ads_df.head())

print("\n=== Dtypes (before) ===")
print("sales_df:\n", sales_df.dtypes)
print("ads_df:\n", ads_df.dtypes)

```

```

=== Raw shapes ===
sales_df: (2184, 6) ads_df: (455, 7)

```

```

=== Sales head ===

```

	Date	Location	Product	Units_Sold	Unit_Price	Revenue
0	2024-01-01	Tbilisi	Espresso	12	3.0	36.0
1	2024-01-01	Tbilisi	Latte	6	4.0	24.0
2	2024-01-01	Tbilisi	Cappuccino	11	4.0	44.0
3	2024-01-01	Tbilisi	Americano	14	4.0	56.0
4	2024-01-01	Tbilisi	Mocha	7	3.5	24.5

=== Ads head ===

	Date	Campaign	Clicks	Impressions	Spend	CTR	CPC
0	2024-01-01	Facebook Ads	25	875	7.69	0.0286	0.31
1	2024-01-01	Google Ads	23	966	4.18	0.0238	0.18
2	2024-01-01	Instagram Ads	15	330	6.48	0.0455	0.43
3	2024-01-01	Billboard	22	1034	9.50	0.0213	0.43
4	2024-01-01	Flyers	18	882	3.19	0.0204	0.18

=== Dtypes (before) ===

sales_df:

```
Date          object
Location       object
Product        object
Units_Sold     int64
Unit_Price     float64
Revenue        float64
```

dtype: object

ads_df:

```
Date          object
Campaign       object
Clicks         int64
Impressions    int64
Spend          float64
CTR            float64
CPC            float64
```

dtype: object

```
[4]: # ----- 3) Basic cleaning -----
# normalize column names
sales_df.columns = [c.strip().replace(" ", "_") for c in sales_df.columns]
ads_df.columns   = [c.strip().replace(" ", "_") for c in ads_df.columns]

# parse dates
for df in (sales_df, ads_df):
    if "Date" in df.columns:
        df["Date"] = pd.to_datetime(df["Date"], errors="coerce")

# coerce numerics
num_cols_sales = ["Units_Sold", "Unit_Price", "Revenue"]
for c in num_cols_sales:
```

```

        if c in sales_df.columns:
            sales_df[c] = pd.to_numeric(sales_df[c], errors="coerce")

num_cols_ads = ["Clicks", "Impressions", "Spend", "CTR", "CPC"]
for c in num_cols_ads:
    if c in ads_df.columns:
        ads_df[c] = pd.to_numeric(ads_df[c], errors="coerce")

```

```

[5]: # ----- 4) Nulls & basic fixes -----
print("\n=== Null counts (after coercion) ===")
print("sales_df:\n", sales_df.isna().sum())
print("ads_df:\n", ads_df.isna().sum())

# Drop rows with missing Date
sales_df = sales_df.dropna(subset=["Date"])
ads_df = ads_df.dropna(subset=["Date"])

# Negative or impossible values -> fix or drop
if "Units_Sold" in sales_df:
    sales_df["Units_Sold"] = sales_df["Units_Sold"].clip(lower=0)

if "Unit_Price" in sales_df:
    sales_df["Unit_Price"] = sales_df["Unit_Price"].clip(lower=0)

if "Revenue" in sales_df:
    missing_rev = sales_df["Revenue"].isna()
    can_compute = missing_rev & sales_df["Units_Sold"].notna() &
    ↪ sales_df["Unit_Price"].notna()
    sales_df.loc[can_compute, "Revenue"] = (sales_df.loc[can_compute,
    ↪ "Units_Sold"] *
                                           sales_df.loc[can_compute,
    ↪ "Unit_Price"]).round(2)
    sales_df["Revenue"] = sales_df["Revenue"].fillna(0).clip(lower=0)

for c in ["Clicks", "Impressions", "Spend"]:
    if c in ads_df:
        ads_df[c] = ads_df[c].fillna(0).clip(lower=0)

# Recompute CTR/CPC if needed
if set(["Clicks", "Impressions"]).issubset(ads_df.columns):
    recompute_ctr_mask = ads_df["CTR"].isna() if "CTR" in ads_df.columns else
    ↪ pd.Series(True, index=ads_df.index)
    ads_df.loc[recompute_ctr_mask, "CTR"] = np.where(ads_df["Impressions"] > 0,
                                                    ads_df["Clicks"] /
    ↪ ads_df["Impressions"], 0)

if set(["Spend", "Clicks"]).issubset(ads_df.columns):

```

```

recompute_cpc_mask = ads_df["CPC"].isna() if "CPC" in ads_df.columns else
↳pd.Series(True, index=ads_df.index)
ads_df.loc[recompute_cpc_mask, "CPC"] = np.where(ads_df["Clicks"] > 0,
ads_df["Spend"] /
↳ads_df["Clicks"], 0)

```

=== Null counts (after coercion) ===

```

sales_df:
  Date      0
Location    0
Product     0
Units_Sold  0
Unit_Price  0
Revenue     0
dtype: int64
ads_df:
  Date      0
Campaign    0
Clicks      0
Impressions 0
Spend       0
CTR         0
CPC         0
dtype: int64

```

```

[6]: # ----- 5) Alignment checks: what can we join on? -----

# 1) What columns overlap?
common_cols = set(sales_df.columns) & set(ads_df.columns)
print("\n=== Common columns between sales_df and ads_df ===")
print(common_cols)

# 2) Quick cardinality overview
print("\n=== Uniques overview ===")
if {"Location", "Product"}.issubset(sales_df.columns):
    print("sales_df uniques -> Locations:", sales_df["Location"].nunique(),
          "| Products:", sales_df["Product"].nunique())
if "Campaign" in ads_df.columns:
    print("ads_df uniques -> Campaigns:", ads_df["Campaign"].nunique())

# 3) Date coverage comparison
sales_dates = set(sales_df["Date"].dropna().unique())
ads_dates   = set(ads_df["Date"].dropna().unique())

only_sales_dates = sales_dates - ads_dates
only_ads_dates   = ads_dates - sales_dates

```

```

both_dates          = sales_dates & ads_dates

print("\n=== Date coverage ===")
print("Sales dates:", len(sales_dates),
      "| Ads dates:", len(ads_dates),
      "| Overlap:", len(both_dates))
print("Dates only in sales_df:", len(only_sales_dates))
print("Dates only in ads_df:", len(only_ads_dates))

# 4) Per-day totals and where data is missing on one side
sales_by_day = (sales_df.groupby("Date", as_index=False)
                .agg(Daily_Revenue=("Revenue", "sum"),
                    Units_Sold=("Units_Sold", "sum")))
ads_by_day = (ads_df.groupby("Date", as_index=False)
              .agg(Ad_Spend=("Spend", "sum"),
                  Ad_Clicks=("Clicks", "sum"),
                  Ad_Impressions=("Impressions", "sum")))

day_merge = sales_by_day.merge(ads_by_day, on="Date", how="outer")

# Flags: which days have which data?
day_merge["has_sales"] = day_merge["Daily_Revenue"].fillna(0) > 0
day_merge["has_ads"]   = day_merge["Ad_Spend"].fillna(0) > 0
day_merge["where"] = np.select(
    [ day_merge["has_sales"] & day_merge["has_ads"],
      day_merge["has_sales"] & ~day_merge["has_ads"],
      ~day_merge["has_sales"] & day_merge["has_ads"] ],
    ["both", "sales_only", "ads_only"],
    default="neither"
)

print("\n=== Day-level availability ===")
print(day_merge["where"].value_counts())

print("\nExamples of days with sales but no ads:")
print(day_merge.loc[day_merge["where"]=="sales_only"].head(5))

print("\nExamples of days with ads but no sales:")
print(day_merge.loc[day_merge["where"]=="ads_only"].head(5))

# 5) Sanity: do we have any negative or impossible values left?
def bad_counts(df, cols):
    out = {}
    for c in cols:
        if c in df.columns:
            out[c] = int((df[c] < 0).sum())
    return out

```

```

print("\n=== Negative-value checks ===")
print("sales_df:", bad_counts(sales_df, ["Units_Sold", "Unit_Price", "Revenue"]))
print("ads_df:", bad_counts(ads_df, ["Clicks", "Impressions", "Spend"]))

# 6) Optional: show a tiny summary of per-campaign completeness (ads side)
if "Campaign" in ads_df.columns:
    camp_days = (ads_df.groupby("Campaign")["Date"]
                  .nunique()
                  .sort_values(ascending=False)
                  .head(10))
    print("\nTop campaigns by active days:")
    print(camp_days)

```

```

=== Common columns between sales_df and ads_df ===
{'Date'}

```

```

=== Uniques overview ===
sales_df uniques -> Locations: 4 | Products: 6
ads_df uniques -> Campaigns: 5

```

```

=== Date coverage ===
Sales dates: 91 | Ads dates: 91 | Overlap: 91
Dates only in sales_df: 0
Dates only in ads_df: 0

```

```

=== Day-level availability ===
where
both      91
Name: count, dtype: int64

```

```

Examples of days with sales but no ads:
Empty DataFrame
Columns: [Date, Daily_Revenue, Units_Sold, Ad_Spend, Ad_Clicks, Ad_Impressions,
has_sales, has_ads, where]
Index: []

```

```

Examples of days with ads but no sales:
Empty DataFrame
Columns: [Date, Daily_Revenue, Units_Sold, Ad_Spend, Ad_Clicks, Ad_Impressions,
has_sales, has_ads, where]
Index: []

```

```

=== Negative-value checks ===
sales_df: {'Units_Sold': 0, 'Unit_Price': 0, 'Revenue': 0}
ads_df: {'Clicks': 0, 'Impressions': 0, 'Spend': 0}

```

Top campaigns by active days:

Campaign	
Billboard	91
Facebook Ads	91
Flyers	91
Google Ads	91
Instagram Ads	91

Name: Date, dtype: int64

```
[7]: # ----- 6) Aggregate daily -----
sales_daily = (sales_df
                .groupby("Date", as_index=False)
                .agg(Daily_Revenue=("Revenue", "sum"),
                    Units_Sold=("Units_Sold", "sum"),
                    Line_Items=("Revenue", "size")))

sales_daily["AOV_proxy"] = np.where(sales_daily["Line_Items"] > 0,
                                    sales_daily["Daily_Revenue"] /
                                    ↪sales_daily["Line_Items"], 0)

ads_daily = (ads_df
             .groupby("Date", as_index=False)
             .agg(Ad_Clicks=("Clicks", "sum"),
                 Ad_Impressions=("Impressions", "sum"),
                 Ad_Spend=("Spend", "sum")))

ads_daily["CTR"] = np.where(ads_daily["Ad_Impressions"] > 0,
                            ads_daily["Ad_Clicks"] /
                            ↪ads_daily["Ad_Impressions"], 0)
ads_daily["CPC"] = np.where(ads_daily["Ad_Clicks"] > 0,
                            ads_daily["Ad_Spend"] / ads_daily["Ad_Clicks"], 0)
ads_daily["CPM"] = np.where(ads_daily["Ad_Impressions"] > 0,
                            (ads_daily["Ad_Spend"] /
                            ↪ads_daily["Ad_Impressions"]) * 1000, 0)

print("\n=== sales_daily preview ===")
print(sales_daily.head())
print("\n=== ads_daily preview ===")
print(ads_daily.head())
```

=== sales_daily preview ===

	Date	Daily_Revenue	Units_Sold	Line_Items	AOV_proxy
0	2024-01-01	748.5	214	24	31.187500
1	2024-01-02	851.5	251	24	35.479167
2	2024-01-03	829.5	254	24	34.562500
3	2024-01-04	826.0	243	24	34.416667

4	2024-01-05	816.0	247	24	34.000000
---	------------	-------	-----	----	-----------

```
=== ads_daily preview ===
```

	Date	Ad_Clicks	Ad_Impressions	Ad_Spend	CTR	CPC \
0	2024-01-01	103	4087	31.04	0.025202	0.301359
1	2024-01-02	115	4579	31.12	0.025115	0.270609
2	2024-01-03	91	2941	30.53	0.030942	0.335495
3	2024-01-04	104	4019	25.68	0.025877	0.246923
4	2024-01-05	102	3992	33.87	0.025551	0.332059

	CPM
0	7.594813
1	6.796244
2	10.380823
3	6.389649
4	8.484469

```
[8]: # ----- 7) Merge & KPIs -----
daily = pd.merge(sales_daily, ads_daily, on="Date", how="left").fillna(0)

daily["ROAS"] = np.where(daily["Ad_Spend"] > 0,
                        daily["Daily_Revenue"] / daily["Ad_Spend"], np.nan)

daily["ROI"] = np.where(daily["Ad_Spend"] > 0,
                        (daily["Daily_Revenue"] - daily["Ad_Spend"]) /
                        ↪daily["Ad_Spend"], np.nan)

daily["ConvRate_proxy"] = np.where(daily["Ad_Clicks"] > 0,
                                   daily["Line_Items"] / daily["Ad_Clicks"], np.
                                   ↪nan)

print("\n=== daily merged preview ===")
print(daily.head())

print("\n=== Summary ranges ===")
print(daily.describe(include="all").T)
```

```
=== daily merged preview ===
```

	Date	Daily_Revenue	Units_Sold	Line_Items	AOV_proxy	Ad_Clicks \
0	2024-01-01	748.5	214	24	31.187500	103
1	2024-01-02	851.5	251	24	35.479167	115
2	2024-01-03	829.5	254	24	34.562500	91
3	2024-01-04	826.0	243	24	34.416667	104
4	2024-01-05	816.0	247	24	34.000000	102

	Ad_Impressions	Ad_Spend	CTR	CPC	CPM	ROAS \
0	4087	31.04	0.025202	0.301359	7.594813	24.114046

1	4579	31.12	0.025115	0.270609	6.796244	27.361825
2	2941	30.53	0.030942	0.335495	10.380823	27.169997
3	4019	25.68	0.025877	0.246923	6.389649	32.165109
4	3992	33.87	0.025551	0.332059	8.484469	24.092117

	ROI	ConvRate_proxy
0	23.114046	0.233010
1	26.361825	0.208696
2	26.169997	0.263736
3	31.165109	0.230769
4	23.092117	0.235294

=== Summary ranges ===

	count	mean	min	\
Date	91	2024-02-15 00:00:00	2024-01-01 00:00:00	
Daily_Revenue	91.0	781.708791	670.0	
Units_Sold	91.0	239.043956	194.0	
Line_Items	91.0	24.0	24.0	
AOV_proxy	91.0	32.5712	27.916667	
Ad_Clicks	91.0	100.065934	79.0	
Ad_Impressions	91.0	3529.241758	2534.0	
Ad_Spend	91.0	29.274945	16.23	
CTR	91.0	0.028642	0.023207	
CPC	91.0	0.292115	0.175567	
CPM	91.0	8.343425	4.508869	
ROAS	91.0	27.864668	17.538393	
ROI	91.0	26.864668	16.538393	
ConvRate_proxy	91.0	0.241985	0.2	

	25%	50%	75%	\
Date	2024-01-23 12:00:00	2024-02-15 00:00:00	2024-03-08 12:00:00	
Daily_Revenue	734.25	779.5	825.5	
Units_Sold	226.0	239.0	251.0	
Line_Items	24.0	24.0	24.0	
AOV_proxy	30.59375	32.479167	34.395833	
Ad_Clicks	93.5	101.0	106.0	
Ad_Impressions	3195.0	3507.0	3837.0	
Ad_Spend	25.175	30.17	32.71	
CTR	0.026192	0.028348	0.03065	
CPC	0.25889	0.286847	0.32543	
CPM	7.221094	8.331432	9.27033	
ROAS	23.596522	26.131814	32.325659	
ROI	22.596522	25.131814	31.325659	
ConvRate_proxy	0.226415	0.237624	0.256692	

	max	std
Date	2024-03-31 00:00:00	NaN
Daily_Revenue	918.5	58.753726

Units_Sold	277.0	16.663274
Line_Items	24.0	0.0
AOV_proxy	38.270833	2.448072
Ad_Clicks	120.0	9.380597
Ad_Impressions	4694.0	481.047014
Ad_Spend	44.93	5.809625
CTR	0.035769	0.002968
CPC	0.417905	0.049155
CPM	12.942092	1.536206
ROAS	48.890943	6.491257
ROI	47.890943	6.491257
ConvRate_proxy	0.303797	0.023258

```
[9]: sales_daily.rename(columns={
    "Daily_Revenue": "Daily Revenue",
    "Units_Sold": "Units Sold",
    "Line_Items": "Line Items",
    "AOV_proxy": "AOV (Proxy)"
}, inplace=True)

ads_daily.rename(columns={
    "Ad_Clicks": "Ad Clicks",
    "Ad_Impressions": "Ad Impressions",
    "Ad_Spend": "Ad Spend",
    "CTR": "Click-Through Rate",
    "CPC": "Cost Per Click",
    "CPM": "Cost Per 1000 Impressions"
}, inplace=True)

daily.rename(columns={
    "Daily_Revenue": "Daily Revenue",
    "Units_Sold": "Units Sold",
    "Line_Items": "Line Items",
    "AOV_proxy": "AOV (Proxy)",
    "Ad_Clicks": "Ad Clicks",
    "Ad_Impressions": "Ad Impressions",
    "Ad_Spend": "Ad Spend",
    "ConvRate_proxy": "Conversion Rate (Proxy)"
}, inplace=True)
```

```
[10]: # ----- 8) Save outputs -----
sales_clean_path = "coffee_sales_clean.csv"
ads_clean_path   = "coffee_ads_clean.csv"
daily_path       = "coffee_daily_kpis.csv"

sales_df.to_csv(sales_clean_path, index=False)
```

```
ads_df.to_csv(ads_clean_path, index=False)
daily.to_csv(daily_path, index=False)

print(f"\nSaved:\n- {sales_clean_path}\n- {ads_clean_path}\n- {daily_path}")
```

Saved:

- coffee_sales_clean.csv
- coffee_ads_clean.csv
- coffee_daily_kpis.csv

[]:

[]:

[]:

[]:

[]:

[]:

[]:

[]:

[]: